

Supplementary Material

Alignment in Multimodal Large Language Models: A Survey

Author information pending confirmation



S1 SCOPE AND READING GUIDE

This independent supplement expands the mechanism and evaluation comparisons in the survey. It contains analytical distinctions and source locations, not new experiments or a matched model leaderboard. The Modality Gap concerns correspondence and use of task-relevant multimodal evidence; the Intent Gap concerns behavior relative to specified goals and constraints. A method can influence both, and a capability error can remain unattributed. Gap labels describe the intervention's intended target, not a claim that it closes the gap.

S2 ARCHITECTURE COMPARISON

Table S1 expands the evidence-access attributes compared in the main text. The relevant cost depends on input resolution, context length, trainable components, and inference budget; qualitative trade-offs should not be read as experimentally matched efficiency rankings.

S3 OPERATIONAL DIAGNOSTIC MATRIX

The contrasts in Table S2 test different hypotheses. Maintain task validity after an edit, hold unrelated inputs fixed, and use a reference appropriate to the edited example. A response change under an arbitrary perturbation is not by itself evidence of grounding.

S4 RELATED-SURVEY COVERAGE LOCATIONS

Table S3 supports the main comparison matrix. Dedicated coverage, distributed discussion, and absence of an organizing dimension are different judgments. A cross does not mean that a paper never mentions a topic. Version-specific locations are used where a reviewed preprint differs from its later publication; these are not claims about every version.

S5 INTERPRETATION AND SCOPE BOUNDARIES

The contribution of the main survey is the joint synthesis of intervention mechanisms, supervision assumptions, diagnostic evidence, and application constraints, not the invention of every axis in the comparison matrix. The comparisons retain foundational vision-language models when they explain an MLLM mechanism and adjacent applications when they expose a transferable constraint. Core MLLM instruction or behavioral alignment is distinguished from image-text representation learning, multimodal entity alignment, and multimodal prediction without an LLM.

Graph augmentation is treated as an architectural mechanism, with application evidence used to explain its consequences. Recommendation and climate cases illustrate supervision and evidence constraints without claiming exhaustive domain coverage. GUI-agent control and multimodal content generation are outside comprehensive coverage. Text-only preference-learning studies are methodological background, not direct demonstrations of multimodal alignment.

TABLE S1
Architectural interventions and the controlled comparisons needed to evaluate them.

Intervention / examples	Selection mechanism	Assumption to inspect	Controlled comparison
Bridge compression: BLIP-2, TokenPacker [1], [2]	Learned query bottleneck or regional-detail injection	A compact token set retains evidence needed downstream	Match decoder and token budget; test small objects and local attributes separately from global answers.
Discrete interfaces: SpeechT5, AnyGPT [3], [4]	Codebook-based representations and token sequences	Discretization preserves task-relevant information across modalities	Separate codebook reconstruction error from reasoning and sequence-length effects.
Explicit structure: MAIL, VideoTree [5], [6]	Entity-mediated exchange or query-adaptive video hierarchy	Extraction and pruning retain decisive entities or events	Perturb or remove selected and unselected evidence; measure missed-event errors and selection cost.
Spatial/layer access: AGE-VLM, Dense Connector [7], [8]	Spatial supervision or multi-layer feature aggregation	Additional features contribute useful detail rather than redundancy	Match image resolution and compute; isolate feature source, fusion, and spatial supervision.
Modality-indexed experts: VLMo, MoT [9], [10]	Modality-conditioned parameter selection with shared interaction	Specialization preserves cross-modal transfer	Report total and active parameters, per-modality loss, and held-out cross-modal tasks.
Adaptive experts: OneLLM, MoIE [11], [12]	Learned routing among projections or experts	Routing uses task-relevant evidence rather than easy modality cues	Audit assignment distributions and evidence interventions separately; balanced utilization is not evidence reliance.
Graph augmentation: LLaVA-SG, MEAformer [13], [14]	Scene relations or multimodal entity structure	Supplied relations are relevant and sufficiently correct	Compare correct, corrupted, and removed relations; distinguish the MLLM mechanism from entity-alignment foundations.

TABLE S2
Diagnostic contrasts for separating evidence failures, constraint violations, and reasoning-capability errors.

Question	Contrast and reference	Supported interpretation	Boundary / control
Is evidence represented?	Probe or compare representations with known modality correspondence; SAIL and MIR [15], [16].	Evidence availability or cross-modal compatibility.	An informative representation need not be used by the decoder.
Is relevant evidence used?	Change a blue mug to red while keeping the color question fixed; update the correct answer.	Consistent, task-appropriate changes support image dependence.	Include irrelevant-region edits and matched natural examples to detect edit artifacts.
Is a claim supported?	Check object claims against annotations, as in POPE and CHAIR [17], [18].	Unsupported-object response rate.	Does not localize the failure to encoding, fusion, or decoding.
Is the action constraint respected?	Keep the scene fixed and contrast “describe without moving” with an authorized action request.	Violations despite correct perception support an Intent Gap attribution.	Check benign utility; universal refusal is not successful task alignment.
Did reasoning fail after perception?	Verify the read values or objects before testing the subsequent inference; MIRAGE [19].	A capability error can be distinguished from incorrect perception.	Do not force the error into either gap without a further evidence-use or constraint test.
Is the reward judge reliable?	Compare pairwise judgments with human-validated labels in VL-RewardBench and Multimodal RewardBench [20], [21].	Agreement on the benchmark’s categories and candidates.	Audit new policy outputs, presentation order, and style confounders separately.
Do gains transfer?	Re-test evidence and constraint contrasts after domain or modality shifts; include uncertainty and benign-utility checks.	Transfer within the documented target conditions.	Aggregate accuracy can hide opposite movements in grounding and behavioral compliance.

TABLE S3
Source locations underlying the related-survey comparison.

Reviewed survey	Coverage locations	Interpretation for the matrix
Caffagni et al. [22], Findings ACL 2024	Sec. 2: model components and training; Sec. 3: capabilities and other modalities; Appendices B–C: training data and evaluation benchmarks.	Objectives are discussed within training, not as a separate objective-family comparison.
Shu et al. [23], Findings EMNLP 2025	Sec. 3: misalignment causes; Sec. 4: mitigation; appendices expand training stages, diagnostics, and benchmarks.	Data, model, training, and evaluation are organizing dimensions; applications are not a dedicated method grouping.
Lu et al. [24], EMNLP 2025	Sec. 3: similarity metrics; Sec. 4: cross-domain evidence; Sec. 5: scale, architectures, objectives, and training/task diversity.	Architecture, objectives, data, training, and transfer applications are discussed, but the organizing focus is representation potential.
Fang et al. [25], publisher PDF	Secs. 3–7: implicit, encoder, behavioral, architectural, and inference-time alignment; Sec. 8: synthetic data and risks; Sec. 9: evaluation.	Already covers representation and behavior, inference-time intervention, data risks, and evaluation. Those topics alone are not exclusive contributions of this survey.
Li and Tang [26], arXiv v2, 11 Oct. 2025	Sec. 2.2: datasets; Sec. 3.4: applications; Sec. 4: fusion methods; Sec. 5: challenges with training-module comparisons.	Applications are explicitly covered. Training and evaluation occur within method/challenge comparisons rather than a separate staged-training and diagnostic framework.
Ho et al. [27], publisher PDF	Sec. 4 groups mitigation by data, architecture, inference, loss optimization, and evaluation.	Substantial lifecycle coverage is acknowledged; application domains are not the central organizing structure.
Yu et al. [28], arXiv v2, 23 Mar. 2025	Sec. 3: preference algorithms and applications; Sec. 4: annotation; Sec. 5: evaluation, including rewards and alignment.	Preference methods, supervision, evaluation, and applications are covered; architectural families are not the organizing focus.

REFERENCES

- [1] J. Li, D. Li, S. Savarese, and S. Hoi, “BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models,” in *International Conference on Machine Learning*. PMLR, 2023, pp. 19730–19742.
- [2] W. Li, Y. Yuan, J. Liu, D. Tang, S. Wang, J. Qin, J. Zhu, and L. Zhang, “TokenPacker: Efficient visual projector for multimodal LLM,” *International Journal of Computer Vision*, vol. 133, no. 10, pp. 6794–6812, 2025.
- [3] J. Ao, R. Wang, L. Zhou, C. Wang, S. Ren, Y. Wu, S. Liu, T. Ko, Q. Li, Y. Zhang *et al.*, “SpeechT5: Unified-modal encoder-decoder pre-training for spoken language processing,” in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2022, pp. 5723–5738.
- [4] J. Zhan, J. Dai, J. Ye, Y. Zhou, D. Zhang, Z. Liu, X. Zhang, R. Yuan, G. Zhang, L. Li *et al.*, “AnyGPT: Unified multimodal LLM with discrete sequence modeling,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024, pp. 9637–9662.
- [5] J. Dong, Q. Zhang, H. Zhou, D. Zha, P. Zheng, and X. Huang, “Modality-aware integration with large language models for knowledge-based visual question answering,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024, pp. 2417–2429.
- [6] Z. Wang, S. Yu, E. Stengel-Esklin, J. Yoon, F. Cheng, G. Bertasius, and M. Bansal, “VideoTree: Adaptive tree-based video representation for LLM reasoning on long videos,” in *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2025, pp. 3272–3283.
- [7] S. Mahajan, H. Le, H. Park, F. Farhadzadeh, M. Hayat, and F. Porikli, “Attention Guided Alignment in Efficient Vision-Language Models,” *arXiv preprint arXiv:2511.17793*, 2025.
- [8] H. Yao, W. Wu, T. Yang, Y. Song, M. Zhang, H. Feng, Y. Sun, Z. Li, W. Ouyang, and J. Wang, “Dense Connector for MLLMs,” in *Advances in Neural Information Processing Systems*, vol. 37, 2024, pp. 33108–33140.
- [9] H. Bao, W. Wang, L. Dong, Q. Liu, O. K. Mohammed, K. Aggarwal, S. Som, S. Piao, and F. Wei, “VLMo: Unified vision-language pre-training with mixture-of-modality-experts,” in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 32897–32912.
- [10] W. Liang, L. Yu, L. Luo, S. Iyer, N. Dong, C. Zhou, G. Ghosh, M. Lewis, W.-t. Yih, L. Zettlemoyer *et al.*, “Mixture-of-transformers: A sparse and scalable architecture for multi-modal foundation models,” *Transactions on Machine Learning Research*, 2025. [Online]. Available: <https://openreview.net/forum?id=Nu6N69i8SB>
- [11] J. Han, K. Gong, Y. Zhang, J. Wang, K. Zhang, D. Lin, Y. Qiao, P. Gao, and X. Yue, “OneLLM: One Framework to Align All Modalities with Language,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 26584–26595.
- [12] D. Wang, S. Wang, Z. Li, Y. Wang, Y. Li, D. Tang, X. Shen, X. Huang, and Z. Wei, “MoIE: Mixture of intra-and inter-modality experts for large vision language models,” *arXiv preprint arXiv:2508.09779*, 2025.
- [13] J. Wang, J. Ju, J. Luan, and Z. Deng, “LLaVA-SG: Leveraging Scene Graphs as Visual Semantic Expression in Vision-Language Models,” in *ICASSP 2025 - 2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2025, pp. 1–5. [Online]. Available: <https://doi.org/10.1109/icassp49660.2025.10887586>
- [14] Z. Chen, J. Chen, W. Zhang, L. Guo, Y. Fang, Y. Huang, Y. Zhang, Y. Geng, J. Z. Pan, W. Song, and H. Chen, “MEAformer: Multi-modal Entity Alignment Transformer for Meta Modality Hybrid,” in *Proceedings of the 31st ACM International Conference on Multimedia*, 2023, pp. 3317–3327. [Online]. Available: <https://doi.org/10.1145/3581783.3611786>
- [15] L. Zhang, Q. Yang, and A. Agrawal, “Assessing and learning alignment of unimodal vision and language models,” in *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2025, pp. 14604–14614.
- [16] Q. Huang, X. Dong, P. Zhang, Y. Zang, Y. Cao, J. Wang, W. Zhang, and N. Yu, “Deciphering Cross-Modal Alignment in Large Vision-Language Models via Modality Integration Rate,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2025, pp. 218–227.
- [17] Y. Li, Y. Du, K. Zhou, J. Wang, X. Zhao, and J.-R. Wen, “Evaluating object hallucination in large vision-language models,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023, pp. 292–305.
- [18] A. Rohrbach, L. A. Hendricks, K. Burns, T. Darrell, and K. Saenko, “Object hallucination in image captioning,” in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018, pp. 4035–4045.
- [19] B. Dong, M. Ni, Z. Huang, G. Yang, W. Zuo, and L. Zhang, “MIRAGE: Assessing hallucination in multimodal reasoning chains of MLLM,” in *Advances in Neural Information Processing Systems*, vol. 38, 2025, pp. 122910–122955.
- [20] L. Li, Y. Wei, Z. Xie, X. Yang, Y. Song, P. Wang, C. An, T. Liu, S. Li, B. Y. Lin, L. Kong, and Q. Liu, “VL-RewardBench: A Challenging Benchmark for Vision-Language Generative Reward Models,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025, pp. 24657–24668. [Online]. Available: https://openaccess.thecvf.com/content/CVPR2025/html/Li_VL-RewardBench_A_Challenging_Benchmark_for_Vision-Language_Generative_Reward_Models_CVPR_2025_paper.html
- [21] M. Yasunaga, L. Zettlemoyer, and M. Ghazvininejad, “Multimodal RewardBench: Holistic Evaluation of Reward Models for Vision Language Models,” *arXiv preprint arXiv:2502.14191*, 2025.
- [22] D. Caffagni, F. Cocchi, L. Barsellotti, N. Moratelli, S. Sarto, L. Baraldi, M. Cornia, and R. Cucchiara, “The Revolution of Multimodal Large Language Models: A Survey,” in *Findings of the Association for Computational Linguistics ACL 2024*, 2024, pp. 13590–13618. [Online]. Available: <https://doi.org/10.18653/v1/2024.findings-acl.807>
- [23] D. Shu, H. Zhao, J. Hu, W. Liu, A. Payani, L. Cheng, and M. Du, “Large Vision-Language Model Alignment and Misalignment: A Survey Through the Lens of Explainability,” in *Findings of the Association for Computational Linguistics: EMNLP 2025*. Association for Computational Linguistics, 2025, pp. 1713–1735. [Online]. Available: <https://aclanthology.org/2025.findings-emnlp.90/>
- [24] J. Lu, H. Wang, Y. Xu, Y. Wang, K. Yang, and Y. Fu, “Representation Potentials of Foundation Models for Multimodal Alignment: A Survey,” in *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2025, pp. 16669–16684. [Online]. Available: <https://aclanthology.org/2025.emnlp-main.843/>
- [25] Y. Fang, X. Pang, Q. Yu, F. Min, X. Cao, P. Tao, and T. Li, “Alignment in Large Vision Language Models: A Survey,” *Information Fusion*, vol. 133, p. 104294, 2026. [Online]. Available: <https://doi.org/10.1016/j.inffus.2026.104294>
- [26] S. Li and H. Tang, “Multimodal Alignment and Fusion: A Survey,” *International Journal of Computer Vision*, vol. 134, 2026. [Online]. Available: <https://doi.org/10.1007/s11263-025-02667-1>
- [27] Z. Y. Ho, S. Liang, and D. Tao, “Understanding Hallucinations in Large Visual and Language Models,” *ACM Computing Surveys*, vol. 58, no. 13, pp. 332:1–332:36, 2026. [Online]. Available: <https://doi.org/10.1145/3811409>
- [28] T. Yu, Y.-F. Zhang, C. Fu, J. Wu, J. Lu, K. Wang, X. Lu, Y. Shen, G. Zhang, D. Song, Y. Yan, T. Xu, Q. Wen, Z. Zhang, Y. Huang, L. Wang, and T. Tan, “Aligning Multimodal LLM with Human Preference: A Survey,” 2025. [Online]. Available: <https://arxiv.org/abs/2503.14504v2>